

AGENT FACTORY

User Manual

Build, test, improve, and operate AI agents
with approved business knowledge.

ENGLISH EDITION

A practical guide for business users

Agent Factory turns approved knowledge into daily work.

Prepare Ontology, Skills, and MCP resources first; Steward only binds resources that already exist.

PURPOSE

Create practical assistants for real business outcomes.

The agent is only as reliable as its instructions and its knowledge.

PREREQUISITES

Ontology, required Skills, and MCP capabilities are created and available before Build.

Validate versions, access, and ownership before starting the Agent conversation.

Ontology provides the knowledge; Agent Factory puts it to work.

01

Ontology

Model, organize, and validate business knowledge.

Reusable, approved knowledge base.

02

Agent Factory

Create agents that use approved knowledge to support business work.

Configured agents that users can run and govern.

03


Agent Steward

Guide the conversation from business need to working agent.

A clear path to create, test, and improve.

Approved knowledge is the shared contract between business intent and agent behavior.


After prerequisites are ready, Build follows six guided steps.



Build selects prepared resources; it does not create the Ontology, Skill, or MCP assets.

The sidebar separates creation, operations, and governance.


Your available options depend on role and workspace configuration.



START HERE


- 01** Open Steward to build or improve
- 02** Open Agents to manage an agent
- 03** Use Control to review activity
- 04** Open Ontology for business knowledge
- 05** Use Settings to configure access

Agents lists published assistants; Teams organizes shared work.



After publishing from the Spec editor, find the Agent in Agents > Published.


Prepare reusable Skills before Build; bind them to the Agent when needed.



BEFORE BUILD

- 01 Review available Skills
- 02 Create or import missing packages
- 03 Confirm version and scope


Prepare MCP capabilities before Build; authorize them before Run.



BEFORE BUILD

- 01 Review available MCP servers
- 02 Add any missing capability
- 03 Validate access before Run

Dashboard summarizes activity; Tasks and Monitor help investigate execution.





The 'Tasks' interface displays a list of recent agent task executions. It includes a search bar, filters for 'Filter by initiator...', 'All Agents', and 'All Status'. The table below lists tasks with columns for Task ID, Name, Initiator, Status, and Duration.

Task ID	Name	Initiator	Status	Duration
TASK_649AE4B07E4	mail-db-ontology-assistant	eben xu	Completed	50s
TASK_816A51956C7	wmc-ontology-mail-assistant	eben xu	Completed	34s
TASK_37053FA8266D	oms-order-alert-agent	eben xu	Completed	1m 10s
TASK_81B3EF199D66	PRD 生成助手	eben xu	Failed	2s
TASK_457CD1D67C47	solution-quotation-agent-3	eben xu	Completed	15s
TASK_B2F0DF369305	solution-quotation-agent-3	eben xu	Completed	1m 30s
TASK_0A6C2540B2ED	solution-quotation-agent-3	eben xu	Completed	50s
TASK_3BF2EA926330	solution-quotation-agent-3	eben xu	Completed	29s
TASK_BFD06916839	solution-quotation-agent-3	eben xu	Completed	32s
TASK_37DAF3A8740F	Copy of inventory-transaction-query-agent	eben xu	Completed	49s
TASK_F947D45C1CE4	Copy of inventory-transaction-query-agent	eben xu	Completed	56s
TASK_BF5D691370B0	Copy of wmc-felo-audit-agent-s69a	eben xu	Completed	54s



Use Dashboard for the overview, Tasks for the execution list, and Monitor for live Agent status.

Approvals and triggers determine when a Run can continue and when it starts.




Human approval can pause a Run. Configure schedules and trigger conditions before starting automated work.

Prepare Ontology and external access before Build.



Create and publish the knowledge base first. During Build, select an existing Ontology; configure Connectors for Run access.


Steward is the conversational workspace for the whole agent lifecycle.



USE STEWARD TO

- 01 Describe a business outcome
- 02 Select approved knowledge
- 03 Test, debug, and improve


Step 1 - start a conversation in Steward.



START THE CONVERSATION

- 01 Choose a suggested prompt
- 02 Enter the request
- 03 Add attachments or actions
- 04 Send the message


Step 2 - describe the outcome, then respond to the first follow-up.



MAKE THE REQUEST CLEAR

- 01 State the user and business outcome
- 02 Answer the scope question from Steward
- 03 Keep the requested work explicit


Step 3 - select an existing Ontology in Steward.



SELECT EXISTING KNOWLEDGE

- 01 Select only the prepared libraries
- 02 Submit the selection to Steward

Step 4 - review evidence and resolve the remaining gaps.



COMPLETE THE BUILD CHECK

- 01 Confirm the Agent goal
- 02 Review runtime evidence
- 03 Resolve current gaps
- 04 Approve generation when ready

Step 5 - review the proposed design before creating it.

1 Agent goal

- Scope
 - Answers business questions across WMS, OMS, and FMS
 - Uses live ontology lookup at run time
 - Reply style: direct answer + short evidence summary

2 Knowledge and boundaries

- Domain
 - Read the user's question
 - Find the right business concepts, relationships, rules, or process details in the selected ontology-backed answer in plain language
 - Add a short "why" summary based on the supporting ontology evidence
 - Say when evidence is incomplete instead of guessing
- Verified readiness
 - Runtime ontology access is prepared for the selected libraries
 - Representative coverage was found for OMS order/shipment concepts and FMS shipment/trip/task/rule concepts
 - The agent is limited to those three ontology domains
- Limits
 - It won't answer beyond WMS, OMS, and FMS
 - It won't invent unsupported facts
 - In unclear cases, it will report uncertainty rather than bluff

3 Generate agent

Generate this agent?

This will create the ontology assistant with live WMS, OMS, and FMS question answering and a short evidence summary in each reply.

☒ Generate agent
Create the agent with the verified workflow above.

☐ Adjust scope
Change domains, answer style, or behavior before generation.

☐ Investigate more
Check more ontology coverage before generating.

Cancel Submit

Send a message...

Agent build draft

Curator Paper

Agent build draft

Curator keeps the readable draft here while Steward explores in the chat. New source-backed lines are highlighted; replaced wording is crossed out.

Agent Goal Working Path Discovery Findings Verification Build Contract Handoff

1. Agent Goal Working

What to build

Build an ontology assistant that answers business questions across WMS, OMS, and FMS using live ontology lookup and includes a short evidence summary with each answer.

Success criteria / boundaries

- Answers business questions across WMS, OMS, and FMS
- Uses live ontology lookup at runtime
- Includes a short evidence summary with each answer
- States uncertainty when ontology evidence is incomplete
- Only uses the selected ontology domains
- Does not invent unsupported facts
- Does not answer outside the selected ontology scope

Needs user: Confirm the verified workflow to generate the agent

2. Working Path Pending

3. Discovery Findings Pending

The selected runtime scope is WMS, OMS, and FMS. A reusable runtime ontology capability is available and has been prepared for those libraries.

4 Agent build draft


What Steward looked for, which sources were used, and how the minimum answer path is forming.

- Listed available ontology libraries.
- Confirmed the user-selected runtime scope as WMS, OMS, and FMS.

REVIEW THE PROPOSAL

- 01 Agent goal
- 02 Knowledge and boundaries
- 03 Generate Agent
- 04 Agent build draft


Step 6 - create and publish the Agent, then verify it in Agents.



VERIFY THE RESULT

- 01 Check the Agent name and purpose
- 02 Confirm its published status
- 03 Open it from Agents to configure or run


Open the Agent detail page and use Execute to start a run.



RUN ENTRY

- 01 Open the published or draft Agent
- 02 Click Execute in the top-right corner
- 03 Continue in the Run workspace


Use a suggested prompt or describe the business task in the message box.



RUN A TASK

- 01 Choose a suggested prompt or type a request
- 02 Review the response and evidence
- 03 Repeat with a clearer request when needed

Configure the Agent from its detail page before sharing it.




WHERE TO CONFIGURE

- 01 Overview - identity and run access
- 02 Config - runtime and behavior settings
- 03 Connectors - authorized external systems
- 04 Skills - reusable task behavior
- 05 Tasks and Schedules - work and timing

Configure, save, test, then publish.

Edit the Spec, then publish it to the Agents module.




SHOW THE SPEC

- 01 Open the Agent card
- 02 Steward shows the Spec on the right
- 03 Ask Steward to modify it

EDIT AND VERSION

- Edit or upload Spec files
- Roll back with History
- Save and test the draft
- Click Publish (top right)
- Find the published Agent in Agents



Use run results to refine the Agent in place.



IMPROVE THE RIGHT LAYER

- Instructions and tone
- Terminology and boundaries
- Skills and tools
- Knowledge selection
- Escalation rules

Use Control to keep Agent activity visible and reviewable.



Use Approval Requests to review blocked work; use History to roll back Spec versions before publishing.

A good agent is a maintained business capability.

Keep the knowledge current. Test before sharing.
Use human review where the work requires it.

QUICK REFERENCE

Ontology -> Steward -> Test -> Improve -> Operate